

```

import tensorflow as tf
import numpy as np
import gym

# Load cartpole environment
env = gym.make('CartPole-v0')

gamma = 0.99
learning_rate = 0.01
state_size = 4
num_actions = 2
hidden_size = 8
total_episodes = 5000 # Set total number of episodes to train agent on.
max_ep = 999
update_frequency = 5
is_visualize = False

def discount_rewards(r):
    """ take 1D float array of rewards and compute discounted reward """
    discounted_r = np.zeros_like(r)
    running_add = 0
    for t in reversed(range(0, r.size)):
        running_add = running_add * gamma + r[t]
        discounted_r[t] = running_add
    return discounted_r

class PolicyNetworks(tf.keras.Model):
    def __init__(self):
        super(PolicyNetworks, self).__init__()
        self.hidden_layer_1 = tf.keras.layers.Dense(hidden_size, activation='relu')
        self.output_layer = tf.keras.layers.Dense(num_actions, activation='softmax')

    def call(self, x):
        H1_output = self.hidden_layer_1(x)
        outputs = self.output_layer(H1_output)

        return outputs

def pg_loss(outputs, actions, rewards):
    indexes = tf.range(0, tf.shape(outputs)[0]) * tf.shape(outputs)[1] + actions
    responsible_outputs = tf.gather(tf.reshape(outputs, [-1]), indexes)

    loss = -tf.reduce_mean(tf.math.log(responsible_outputs) * rewards)

    return loss

def train_step(model, states, actions, rewards):
    with tf.GradientTape() as tape:
        outputs = model(states)
        loss = pg_loss(outputs, actions, rewards)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))

# Declare Policy Gradient Networks
PG_model = PolicyNetworks()

i = 0
total_reward = []
total_length = []

# train start
while i < total_episodes:
    s = env.reset()
    running_reward = 0
    ep_history = []

    for j in range(max_ep):
        if is_visualize == True:
            env.render()
        # Probabilistically pick an action given our network outputs.
        s = np.expand_dims(s, 0)
        a_dist = PG_model(s).numpy()
        a = np.random.choice(a_dist[0], p=a_dist[0])
        a = np.argmax(a_dist == a)

```

Saved successfully!



ing_rate)

```

s1, r, d, _ = env.step(a) # Get reward and next state
ep_history.append([s, a, r, s1])
s = s1
running_reward += r

if d == True:
    ep_history = np.array(ep_history)
    ep_history[:, 2] = discount_rewards(ep_history[:, 2])

    # Make state list to numpy array
    np_states = np.array(ep_history[0, 0])
    for idx in range(1, ep_history[:, 0].size):
        np_states = np.append(np_states, ep_history[idx, 0], axis=0)

    # Update the network parameter
    if i % update_frequency == 0 and i != 0:
        train_step(PG_model, np_states, ep_history[:, 1], ep_history[:, 2])

    total_reward.append(running_reward)
    total_length.append(j)
    break

# Print last 100 episode's mean score
if i % 100 == 0:
    print(np.mean(total_reward[-100:]))
i += 1

<ipython-input-53-bba22c6d4444>:84: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tupl
ep_history = np.array(ep_history)
14.0
28.36
32.97
40.33
48.04
37.38
42.99
67.76
62.87
36.61
36.32
36.97
56.17
100.37
113.89
95.48
64.73
96.71
64.36
158.5
182.78
183.9
192.02
190.56
161.07
56.08
71.56
81.91
94.96
92.61
96.14
159.24
197.09
195.93
196.8
189.17
185.5
175.52
157.13
186.29
195.07
190.3
190.25
187.21
162.4
175.18
177.29

```

Saved successfully!



